

BGP Simulator

Protocol Processing
Spring 2017

Group 8
Niko Hellgren

Contents

- Technical specification and limitations
- Program architecture
- BGP session initialization
- Routing engine
- Trust implementation

Technical specification

- Implemented in Java (built-in libraries + GraphStream graph plotting library for visualization purposes)
- Strongly threaded
- Inter-router data transmission (L1 & L2) using Java's Stream classes
- L3 implementation based on IPv4...
 - ...Although not using all the available fields
 - No packet fragmentation support
 - No different types of service
 - No differentiation between protocols

Technical specification

- BGP router data transmission using BGP-4
- For necessary but not essential features, a globally available state is present
 - DNS-like feature for resolving router addresses
 - PKI-like feature for distributing public keys used in trust voting
 - Static functionality to start the connecting process of two routers
- To simulate traffic in the network, clients with pinging ability attached to the routers

Technical limitations

- Emphasis of the simulation is in BGP
 - The notion of Autonomous System (AS) is absorbed into routers, and each AS consists of only one router and its connected users
- TCP is not implemented, BGP messages are transmitted as IP packet payload
 - Attacks and vulnerabilities based on TCP features are not present

routing

SubnetNode

```
# subnet: Subnet
# parent: SubnetNode
# children: List<SubnetNode>
- firstHop: int
- pathLength: int
+ changePath(int firstHop, int pathLength)
+ addChild()
+ delete()
```

RoutingEngine

```
- subnetRootNode : SubnetNode
- trustProvider : TrustProvider
+ decidePath(Address target)
+ getBestMatchingSubnetNode(Address target)
+ getSubnetsBehind(int routerId)
+ handleUpdateMessage(UpdateMessage um)
```

trust

<<Interface>>
TrustProvider

```
+ getTrustFor(int routerId)
```

TrustEngine

```
- kp: KeyPair
- directTrustWeight: double
- votedTrustValues: Map<Integer, Byte>
- directTrustValues: Map<Integer, Byte>
+ getTrustFor(int routerId)
+ handleTrustMessage(TrustMessage tm)
+ changeDirectTrust(int otherId, int delta)
+ getPublicKey()
```

BGPRouter

```
+ id : int
+ subnet : Subnet
- routingEngine: RoutingEngine
- trustEngine : TrustEngine
- connections: Map<Integer, ASConnection>
- packetReceivers: Map<Address, PacketReceiver>
+ routePacket(byte[] packet)
- sendToNeighbour(int id, byte[] packet)
+ reserveAddress()
+ getPublicKey()
+ shutdown()
```

<<Interface>>
PacketRouter

```
+ routePacket(byte[] packet)
```

network

ASConnection

```
- ownAddress: Address
- neighbourId: int
- adapter: InterRouterInterface
- fsm: StateMachine
- keepaliveChecking: TimerTask
- keepaliveSending: TimerTask
+ start()
+ sendOpenMessage()
+ handleOpenMessage()
+ raiseKeepaliveFlag()
+ raiseNotification()
+ closeConnection()
```

InterRouterInterface

```
- processingThread: Thread
- in: PipedInputStream
- out: PipedOutputStream
- conn: ASConnection
+ sendData(byte[] packet)
+ connectNeighbour(InterRouterInterface other)
+ run()
+ close()
```

Program architecture

- `bgp.simulation.SimulatorState`
 - Global state; contains router and client registry, DNS-like functionality to map addresses to routers, and PKI-functionality for trust voting
- `bgp.core.BGPRouter`
 - Base class for routers
 - Each instance is given a unique ID and a subnet
 - Provides a DHCP-like functionality to connected clients
 - Contains.
 - A list of `ASConnection` objects representing the connections to other routers
 - `RoutingEngine` responsible for routing decisions and routing table
 - `TrustEngine` responsible for calculating and voting on trust

Program architecture

- `bgp.core.ASConnection`
 - Represents a connection to another BGPRouter
 - Contains the finite state machine, and timers for checking and sending KEEPALIVE messages
 - Handles the connection initialization process
 - Generates new OPEN, KEEPALIVE and NOTIFICATION messages
- `bgp.core.network.InterRouterInterface`
 - Provides the L1/L2 functionalities for ASConnection
 - Uses Java's built-in `PipedInputStream` and `PipedOutputStream` classes
 - Transfers data as a byte stream, prepended by the length of the packet
 - Passes received packets to BGPRouter for routing

Program architecture

- `bgp.core.routing.RoutingEngine`
 - Handles UPDATE messages, decides the routing paths and keeps up the routing table
- `bgp.core.trust.TrustEngine`
 - Calculates trust values for neighbours, constructs and responds to TRUST messages
- `bgp.client.BGPClient`
 - Represents the clients attached to the network for testing purposes
 - Connected to a single BGPRouter, receives an IP address from it
 - Capable of responding to ping requests

Program architecture

- `bgp.utils.Address`
 - Address and its subclass Subnet represent the IP addresses and subnets, can be converted to byte or long representation
- `bgp.utils.PacketEngine`
 - Functionalities regarding the construction, modification and validation of IPv4 packets

BGP session initialization

1. Create a new ASConnection
2. Connect the InterRouterInterfaces' streams
3. Set the FSM to state CONNECT
4. Send an OPEN message, retry until a KEEPALIVE is received or 10 retries have been done
5. Wait (independently of the previous steps) for OPEN message. Once received, start sending and checking KEEPALIVE messages
6. After KEEPALIVE is received, change state to ESTABLISHED
7. Send own routing table information to the other party, start using the connection for packet routing

Routing engine – Decision process

Validate the received UPDATE message

FOR EACH withdrawn route

IF exactly matching subnet in routing table is not found

Do nothing

ELSE IF the first hop on the UPDATE message matches the first hop of current best path

Remove the routing information for the subnet

Store the subnet to be sent to other neighbours

ELSE

Store the information about own best path to revoked subnet to be sent to the revoking peer after processing

Routing engine – Decision process

FOR EACH NLRI row

IF no exact subnet match is found from routing table

Create a new subnet node to routing table with sending router as first hop

ELSE IF local preference values of existing and new routes' first hops differ

Store the preferred route to routing table

ELSE IF trust-scaled path lengths of existing and new routes' differ

Store the shorter route to routing table

IF selected path was changed

Store new routing information to be sent to other neighbours

Routing engine – Decision process

Modify the UPDATE message:

- Add own ID to AS_PATH

- Remove revoked routes that caused no changes to routing table

- Remove NLRI entries that caused no changes to routing table

FOR EACH neighbour not in the AS_PATH

- Set NEXT_HOP to the address of the interface used to communicate with neighbour

- Send the UPDATE message to neighbour

IF withdrawn routes that did not modify routing table exist

- Send UPDATE message containing NLRI for these subnets to the revoking neighbour

Trust implementation

- Each router holds two trust values for each neighbour: inherent trust and voted trust
- Stored as values in range -128..127 for easier transmission, scaled to 0..1 when used for path cost calculations
- *Normalized routing criterion with trust rate* calculated similarly to He's solution¹:
$$scaled\ cost = \frac{initial\ cost}{trust\ rate}$$
- Trust rate is the weighted average of inherent and voted trust
- Voted trust is the average of received votes on the neighbour
- Inherent trust is specified by operator and neighbour behaviour

1. L. He, "A novel scheme on building a trusted IP routing infrastructure", 2006

Trust implementation – Voting

- Voting is done via specific TRUST messages
 - As no BGP session to the voter is established, the messages are not BGP messages per se, but should be considered as ordinary UDP datagrams
- Messages consist of a flag specifying if they are vote requests or responses, voted ID and target ID. If the vote is a response, the message also contains the vote (encrypted using 1024-bit RSA), and a signature. Public keys are provided by the global state.

Trust implementation – Voting

- Voting is triggered by UPDATE messages with more than one host in the AS_PATH
 - A vote request is sent to the second-order neighbour in the AS_PATH list and a token is stored to ensure only one vote per request is accepted
 - The second-order neighbour responds by taking the trust value, padding the single-byte vote with random bits to avoid brute-force attacks, encrypts the vote using requester's public key, signs the encrypted vote with own private key and sends the response back
 - The original requester checks that the token has not yet been used, validates the signature and decrypts the payload
 - Extracted trust value is then calculated to the averaged voted trust